

GeoAgentBench: Benchmarking LLM Agents for Environmental Geospatial Analysis Across Eight Foundation Models

Gabriel Diaz-Ireland^{1,2} Diego Prieto-Herráez² Javier Velázquez²

Mario García Peces¹ Guillermo Perez¹

¹Darwins Geospatial S.L. ²Universidad Católica de Ávila

Abstract. Environmental scientists spend disproportionate effort on data wrangling rather than analysis, yet no benchmark exists to evaluate AI agents that automate environmental geospatial workflows through structured tool calling against real APIs. We introduce **GeoAgentBench**, a benchmark for environmental analysis agents comprising 93 tasks across 18 categories—including municipality-level analysis, multi-turn conversation, spatial reasoning, cross-indicator synthesis, error handling and recovery, ranking, comparison, multilingual understanding, habitat analysis, deep-dive profiling, temporal change, and task rejection—evaluated against a production geospatial API serving three environmental indicators across Spain and Portugal using twelve domain-specific tools. We evaluate eight large language models (GLM-5, DeepSeek V3.2, GPT-OSS-120B, Gemini 2.5 Pro, Qwen3-235B, Llama 4 Scout, Llama 4 Maverick, and Claude Sonnet 4) across Vertex AI and Anthropic platforms, finding that: (1) the best models achieve 58.1% accuracy (GLM-5 and Claude Sonnet 4), revealing that environmental geospatial tool orchestration remains an open challenge; (2) cost-accuracy trade-offs vary by two orders of magnitude, with DeepSeek V3.2 achieving 52.7% at \$0.008/case while Claude Sonnet 4 costs $11\times$ more for the highest accuracy; (3) comparison tasks remain universally unsolved (0% across all models), exposing systematic reasoning limitations; and (4) structured tool calling against a real API provides a more discriminative evaluation signal than synthetic benchmarks. We demonstrate domain and geographic extensibility by incorporating BigEarthNet V2 land cover data for Portugal as a third indicator alongside Spanish CO2 suitability and gully erosion, enabling cross-country evaluation across three data domains. GeoAgentBench, evaluation code, and the MLOps pipeline are publicly available.¹

Keywords: Geospatial AI, LLM Agents, Tool Calling, Benchmark, Environmental Analysis

1 Introduction

Environmental monitoring at scale demands multi-temporal geospatial analysis: computing vegetation indices, detecting land cover change, assessing erosion risk, and generating statistical summaries for regulatory compliance. Practitioners routinely spend the majority of their effort on data wrangling—discovering datasets, reprojecting coordinate systems, and debugging code against remote sensing APIs [Zhang et al. \[2023\]](#), [Li and Ning \[2023\]](#). This expertise barrier limits the adoption of geospatial methods by domain scientists who understand environmental systems but lack GIS programming skills.

Recent advances in large language models (LLMs) have produced AI agent systems that translate natural language into executable geospatial operations. Architectures range from LLMs with tool pools [Zhang et al. \[2023\]](#), [Akinboyewa et al. \[2024\]](#) to fine-tuned specialists [Zhang et al. \[2024, 2025\]](#) to multi-agent systems achieving 85–97% accuracy [Luo et al. \[2025\]](#), [Lee et al. \[2025\]](#). Major industry players are investing heavily: Google’s Earth AI integrates Gemini with AlphaEarth Foundations [Google Research \[2025\]](#),

the Planet–Anthropic partnership applies Claude to daily satellite imagery [Planet Labs and Anthropic \[2025\]](#), and NASA/IBM’s Prithvi-EO-2.0 provides pre-trained vision transformers [Szwarcman et al. \[2024\]](#).

However, a critical evaluation gap persists. Existing agent benchmarks target general GIS tasks (GeoBenchX [Krechetova et al. \[2025\]](#), ThinkGeo [Shabbir et al. \[2025\]](#)) rather than environmental science workflows. Environmental knowledge benchmarks (EnviroExam [Huang et al. \[2024\]](#)) evaluate knowledge, not agent behaviour. Code-generation benchmarks (UnivEARTH [Kao et al. \[2025\]](#)) do not reflect the structured API interaction that production systems employ—58% of LLM-generated Earth Engine code fails to execute. At the foundation model level, GEO-Bench [Lacoste et al. \[2023\]](#) and GEO-Bench-2 [Simumba et al. \[2025\]](#) evaluate pixel-level vision models, not the agent-level tool orchestration that determines real-world usability.

This paper makes four contributions:

1. **GeoAgentBench**—the first benchmark for environmental analysis agents using structured tool calling against

a production geospatial API, comprising 93 tasks across 18 categories with twelve domain-specific tools.

2. **Eight-model evaluation**—the first systematic comparison of eight LLMs on environmental geospatial tool orchestration under identical conditions across Vertex AI and Anthropic platforms.
3. **Cost-efficiency analysis**—quantitative cost-accuracy trade-off mapping revealing two orders of magnitude variation.
4. **Reproducible MLOps pipeline**—a fully automated Cloud Build pipeline producing structured results with full conversation traces.

2 Related Work

2.1 Geospatial AI Agent Architectures

The literature reveals four paradigms. The **LLM + Tool Pool** approach pairs general-purpose LLMs with callable tools: GeoGPT Zhang et al. [2023] uses GPT-3.5-turbo with LangChain, LLM-Geo Li and Ning [2023] uses GPT-4 for DAG-based decomposition, and GIS Copilot Akinboyewa et al. [2024] integrates LLMs into QGIS. **Fine-tuned specialists** sacrifice generality for expertise: GTChain Zhang et al. [2024] fine-tuned LLaMA-2-7B achieving 32.5% higher accuracy than GPT-4; EnvGPT Zhang et al. [2025] fine-tuned an 8B model on 100M environmental tokens, rivalling GPT-4o-mini. **Multi-agent systems** decompose tasks: GeoJSON Agents Luo et al. [2025] achieved 97% with GPT-4o Planner-Worker architecture vs 49% single-agent; GeoLLM-Squad Lee et al. [2025] gained 17% from specialised sub-agents. **Cost-efficient hybrids**: Geo-OLM Stamoulis et al. [2025] enables sub-7B models to perform within 10% of GPT-4o at 100× lower cost.

2.2 Benchmarks and Evaluation

Agent tool-use is evaluated by GeoBenchX Krechetova et al. [2025], ThinkGeo Shabbir et al. [2025] (486 remote sensing tasks), and UnivEARTH Kao et al. [2025] (only 33% accuracy on Earth Engine code generation, 58% code failure rate). GeoBenchX is the closest comparable work: 202 tasks across five categories (merge-visualise, process-merge-visualise, spatial operations, heatmaps/contour lines, and control questions), evaluated via 24 general-purpose GIS tools (data loading, filtering, spatial joins, choropleth rendering) with an LLM-as-judge scoring protocol (3-point scale, three independent judges achieving 88–96% agreement with human annotations). Top performers are o4-mini and Claude 3.5 Sonnet. Crucially, GeoBenchX tools are low-level primitives operating on local files—, , , —whereas GeoAgentBench tools are high-level domain operations against a production cloud API.

To bridge the gap between general-purpose GIS benchmarks and domain-specific environmental analysis, we integrate four tools from GeoBenchX Krechetova et al. [2025] into GeoAgentBench’s toolset: (proximity queries), (spatial predicates), (point geometry extraction), and (explicit task rejection to prevent hallucination). These tools are adapted from GeoBenchX’s local-file paradigm to operate against GeoAgentBench’s cloud API and administrative boundary system. We additionally introduce 10 benchmark cases inspired by GeoBenchX’s task patterns—buffer-based proximity analysis, spatial joins, and centroid calculations—applied to Spanish environmental indicators. This creates a direct comparison pathway: models achieving 58% on GeoBenchX’s general GIS tasks can be evaluated on equivalent spatial operations within our domain-specific context, isolating the performance impact of domain complexity from tool-use capability.

Map reasoning: MapEval Dihan et al. [2025] (no model exceeds 67%), MapQA Arnold et al. [2025] (humans 91%, models <50%). Spatial reasoning: GPSBench Truong et al. [2026], SpatiaLab Wasi et al. [2026] (best model 55% vs human 88%). Foundation models: GEO-Bench Lacoste et al. [2023] and GEO-Bench-2 Simumba et al. [2025] evaluate *pixel-level* classification/segmentation across 19 datasets; PANGAEA Marsocci et al. [2024] finds foundation models do not consistently outperform supervised baselines.

Key distinction: GEO-Bench evaluates pixel-level model accuracy (*can this model classify a satellite tile?*); GeoAgentBench evaluates agent-level tool orchestration (*can this agent reason about which tools to call?*). The two are complementary.

2.3 Environmental Science AI

EnvGPT Zhang et al. [2025] shows focused fine-tuning on environmental tokens rivals GPT-4o-mini. Google’s Earth AI Google Research [2025] achieves 82% on Q&A (vs 50% baseline Gemini). Pan et al. Pan et al. [2025] demonstrate MCP-grounded air quality monitoring at 4.78/5.0 factual accuracy. Spain’s RD 214/2025 Kingdom of Spain [2025] mandates carbon footprint reporting for ~4,000 organisations, creating demand for automated multi-temporal geospatial analysis.

2.4 The Missing Benchmark

No existing benchmark evaluates environmental analysis agents performing structured tool calling against a real geospatial API with multi-turn conversation, multilingual inputs, and cross-indicator synthesis. GeoAgentBench fills this gap.

Table 1: GeoAgentBench v5 task categories (93 tasks, 18 categories).

Category	N	Tests
Tool selection	21	Correct tool choice for varied queries
Cross-indicator	8	CO ₂ × erosion × land cover synthesis
Deep dive	6	Multi-tool environmental profiling
Interpretation	7	Reasoning over analysis results
Error handling	6	Non-existent entities, invalid indicators
Habitat analysis	7	BigEarthNet V2 land cover (Portugal)
Language	6	Galician, Basque inputs
Municipality	4	Municipality-level analysis, disambiguation
Memory (multi-turn)	6	Follow-up questions, context retention
Spatial reasoning	4	Coastal, meseta, community knowledge
Multi-muni ranking	3	Multi-municipality comparisons
Threshold	3	Numeric threshold queries
Temporal change	1	Cross-country temporal context
Error recovery	3	Graceful fallback when data unavailable
Ranking	2	Top-N and bottom-N province queries
Comparison	2	Close-value and directional comparisons
Single analysis	2	Basic single-province analysis
Province aggregation	2	Aggregate statistics across provinces

3 GeoAgentBench

3.1 Benchmark Design

GeoAgentBench comprises 93 tasks organised into 18 categories (Table 1), spanning easy (19%), medium (45%), and hard (36%) difficulty levels.

Each task specifies: a natural language query (optionally with multi-turn history), expected tool calls, must-contain/must-not-contain strings, maximum rounds, cost budget, and domain-expert ground truth. Error handling tasks are deliberately unsolvable—requesting analysis on non-existent municipalities, fabricating statistics from display-only layers, or citing indicators that do not exist—testing the agent’s ability to decline gracefully rather than hallucinate.

The benchmark covers three environmental indicators across two countries:

- **CO₂ absorption suitability** (categorical, Spain): legislative pre-screening encoding five MITECO spatial criteria [Kingdom of Spain \[2025\]](#), producing three

GeoAgentBench — System Architecture




Figure 1: GeoAgentBench system architecture. The LLM agent interacts with a production API serving COG data; the eval harness logs all tool calls and scores responses.

classes (Not Eligible, Eligible with Conditions, Eligible). Served as a Cloud Optimized GeoTIFF (COG).

- **Gully erosion probability** (continuous, Europe): machine learning prediction from LUCAS 2022 soil survey (0–100%). Served as a COG.
- **BigEarthNet V2 land cover** (categorical, Portugal): 7-class land cover distribution derived from 75k+ labeled Sentinel-2 patches [Clasen et al. \[2024\]](#), aggregated by Portuguese district. Served as pre-computed JSON statistics.

Seven additional layers are available for visual display but not statistical analysis, creating natural hallucination traps.

3.2 Agent Architecture

Figure 1 shows the system architecture. The agent operates in a ReAct-style loop [Yao et al. \[2023\]](#): it receives a query, reasons about tools, executes calls, observes results, and generates a response. The agent has access to twelve tools (Table 2).

The design choice of structured tool calling over code generation is motivated by empirical evidence: GeoJSON Agents report 85–97% accuracy with function calling [Luo et al. \[2025\]](#), while UnivEARTH reports only 33% when agents generate executable code [Kao et al. \[2025\]](#).

Multi-backend inference client. Fair cross-model comparison requires normalizing heterogeneous API surfaces into a common interface. Our inference client routes models through two backends: a native Anthropic client for Claude, and a LiteLLM-based client for Vertex AI Model-as-a-Service (MaaS) endpoints that expose an OpenAI-compatible API. The MaaS backend maps publisher-prefixed model identifiers (e.g.,) to Vertex AI endpoints and normalizes response formats—including tool-call representations—

Table 2: Agent tools. The first eight are domain-specific; the last four are adapted from GeoBenchX Krechetova et al. [2025].

Tool	Purpose
	Resolve province $\rightarrow$ boundary geometry
	Resolve municipality (+ province hint)
	Zonal statistics for indicator in AOI
	Multi-indicator analysis in single call
	Compare two areas on same indicator
	Rank provinces by indicator value
	Enable/disable map layer visibility
	Bar/stacked-bar chart from data
	Buffer geometry by distance (km)
	Select features by spatial predicate
	Extract centroid coordinates
	Decline unsolvable tasks

into the shared schema consumed by the eval harness. Notably, Llama 4 models emit tool calls as Python-style function invocations () rather than JSON; a dedicated AST-based parser extracts these into structured tool calls, enabling Scout to achieve 5.4% accuracy that would otherwise register as 0%. Unattended execution of 744 model–task pairs requires robust failure handling: each inference call uses a layered timeout architecture combining daemon-thread hard timeouts with exponential-backoff retry on rate-limit errors.

3.3 Evaluation Protocol

Each task is scored by eight check types (Table 3 and Figure 2). A task *passes* only when **all** applicable checks pass (strict binary). We additionally compute partial-credit metrics for finer-grained analysis.

Continuous metrics. Beyond binary pass/fail, we report: *check score*—fraction of individual checks passed per case (partial credit); *keyword coverage*—fraction of terms found; *tool F1*—set-based F1 between expected and actual tools; and *cost utilisation*—ratio of actual to budgeted cost.

Error classification. When a task fails, the first failing check (in priority order: `tool_missing` $\rightarrow$ `chart_missing` $\rightarrow$ `wrong_data` $\rightarrow$ `rounds_exceeded` $\rightarrow$ `cost_exceeded` $\rightarrow$ `keyword_missing` $\rightarrow$ `forbidden_keyword`) determines the error category. This priority reflects diagnostic value: tool selection failures are the most informative for understanding agent reasoning.

Design rationale. Tool and action checks use *recall* as the gate rather than F1. An agent that calls all expected tools plus an exploratory extra tool should not be penalised—the key question is whether it performed the required operations. F1 is tracked as a separate metric for efficiency analysis.

Table 3: Scoring checks. Each task specifies which checks apply; a task passes only when every applicable check passes.

Check	Pass condition
	Every expected tool was called (recall = 1.0). Extra tools do not cause failure. ^a
	Every expected UI action was generated (recall = 1.0).
	Each required keyword appears in the answer (case-insensitive substring).
	No forbidden keyword appears in the answer. ^a Tool F1
	For each ground-truth entry, the label is found in the answer and the nearest percentage value is within tolerance. ^b
	At least one chart URL was produced (only checked when is an expected tool).
	Agent loop iterations $\leq$ budget.
	Estimated token cost $\leq$ budget.

(precision, recall, F1 on unique tool sets) is computed as a continuous metric but is not used as a gate. ^bA 120-character window after the label is searched for the first pattern; tolerance defaults to ± 2 pp.

3.4 MLOps Pipeline

Experiments are defined as YAML configuration files. A Cloud Build pipeline builds a Docker image, deploys a Cloud Run Job executing all experiments sequentially, and generates structured outputs: (per-case with full answer, conversation traces, tool inputs/outputs), (frozen config + environment), and a batch-level with leaderboard, per-experiment metrics, and per-case Q&A. All results are grouped under a run-level GCS prefix for isolation and traceability.

4 Experiments and Results

4.1 Experimental Setup

We evaluate eight LLMs on GeoAgentBench v5 (93 tasks, 18 categories) across Vertex AI Model-as-a-Service and the Anthropic native API (Table 4). All experiments use the v3 system prompt (zero-shot), , , and . Experiments run sequentially inside a single Cloud Run Job to ensure identical infrastructure conditions.

4.2 Overall Results

Table 5 and Figure 3 present the main results. GLM-5 and Claude Sonnet 4 share the top position at 58.1%, followed by DeepSeek V3.2 (52.7%) and Qwen3-235B (47.3%). Three models exceed 50%, yet no model surpasses 60%, confirming that environmental geospatial tool orchestration remains a significant challenge even for frontier models. The top five models cluster between 40–58%, with a sharp drop to




Figure 2: Scoring pipeline. Each case is evaluated by up to eight checks; binary pass requires all to pass. Partial-credit metrics (check score, tool F1, keyword coverage) are always computed.

Table 4: Models evaluated. Seven via Vertex AI; Claude via native Anthropic API.

Model	Provider	Params	Access
GLM-5	Zhipu AI	—	Vertex MaaS
DeepSeek V3.2	DeepSeek AI	671B MoE	Vertex MaaS
Qwen3-235B	Alibaba	235B MoE	Vertex MaaS
GPT-OSS-120B	OpenAI	120B	Vertex MaaS
Gemini 2.5 Pro	Google	—	Vertex native
Llama 4 Scout	Meta	109B MoE	Vertex MaaS
Llama 4 Maverick	Meta	400B MoE	Vertex MaaS
Claude Sonnet 4	Anthropic	—	Anthropic API

Llama 4 Scout (5.4%). Llama 4 Maverick returned 0% due to persistent Vertex AI infrastructure errors (HTTP 500) across all 93 cases; we report this as an availability finding rather than a model capability result.

Claude Sonnet 4—the only model accessed via a native provider API rather than Vertex AI MaaS—ties with GLM-5 for first place with notable strengths on municipality analysis, ranking, threshold tasks, and BigEarthNet land cover cases (83% on the 24 BigEarthNet-related cases). However, it is the most expensive model at \$0.087/case, making it 3× costlier than GLM-5.

Partial credit tells a different story. Even failing cases often demonstrate partial competence: correct tool selection but missing a keyword, or correct keywords but an extra unnecessary tool call. This gap between binary accuracy and partial-credit scores (Figure 4) suggests that many failures are “near-misses” addressable through prompt engineering.




Figure 3: GeoAgentBench v5 leaderboard. Accuracy (%) with total cost in parentheses. Dashed line at 50%.

Table 5: Leaderboard: 8 models on 93 GeoAgentBench v5 tasks. Cost is total for all 93 cases.

#	Model	Acc.	Cost	\$/case
1	GLM-5	58.1%	\$2.54	.027
1	Claude Sonnet 4	58.1%	\$8.13	.087
3	DeepSeek V3.2	52.7%	\$0.79	.008
4	Qwen3-235B	47.3%	\$0.50	.005
5	Gemini 2.5 Pro	39.8%	\$2.99	.032
5	GPT-OSS-120B	39.8%	\$4.77	.051
7	Llama 4 Scout	5.4%	\$0.04	.000
8	Llama 4 Maverick	0.0%	—	—

4.3 Cost-Accuracy Analysis

Cost varies by an order of magnitude (Figure 5). Qwen3-235B offers the best cost-efficiency at \$0.005/case with 47.3% accuracy, while Claude Sonnet 4 costs 17× more (\$0.087/case) for the joint-highest accuracy (58.1%). GLM-5 matches Claude’s accuracy (58.1%) at a moderate \$0.027/case. The Pareto frontier is dominated by Qwen3 (cheapest competitive option), DeepSeek (best accuracy-to-cost ratio at 52.7%), and GLM-5 (highest accuracy at lowest cost among top models).

Table 6 shows projected costs for a production deployment. A service processing 1,000 queries/day would pay \$5–8/day with Qwen3 or DeepSeek (47–53% accuracy) vs \$87/day with Claude Sonnet 4 for the highest accuracy (58.1%)—a compelling case for empirical benchmarking over provider reputation. This finding echoes Geo-OLM’s demonstration [Stamoulis et al. \[2025\]](#) that cost-efficient models can match frontier performance at dramatically lower cost.

Figure 6 further illustrates this relationship, showing that token consumption does not predict accuracy—models that generate more tokens do not necessarily achieve higher scores.




Figure 4: Binary accuracy vs partial-credit check score. The gap indicates many “near-miss” failures.




Figure 5: Cost-accuracy trade-off. Bubble size proportional to total tokens. Qwen3-235B and DeepSeek V3.2 dominate the Pareto frontier.

4.4 Failure Analysis

4.4.1 Universally Hard Tasks

One category achieved 0% accuracy across all completed models (Figure 7), revealing a systematic limitation:

Comparison (0/2 for all models): close-value comparisons (e.g., Murcia 68.5% vs Córdoba 68.4%) caused all models to fabricate differences rather than reporting near-identical values. **Deep dive** (0–50%): multi-tool environmental profiling requiring sequential tool calls proved difficult, with most models solving at most two of six tasks. **Error recovery** (0–33%): graceful fallback when data is unavailable was solved inconsistently across models.

Conversely, **error handling** was the easiest category: Qwen3-235B achieved 100% (6/6), and GLM-5, Gemini, and GPT-OSS each scored 67%. Models can recognise invalid inputs but struggle with multi-step analytical reasoning.

Table 6: Projected daily cost at 1,000 queries/day.

Model	Accuracy	Daily Cost	Monthly Cost
Qwen3-235B	47.3%	\$5	\$150
DeepSeek V3.2	52.7%	\$8	\$240
GLM-5	58.1%	\$27	\$810
Gemini 2.5 Pro	39.8%	\$32	\$960
GPT-OSS-120B	39.8%	\$51	\$1,530
Claude Sonnet 4	58.1%	\$87	\$2,610




Figure 6: Token consumption vs accuracy. Token volume is a poor predictor of performance.

4.4.2 Model-Specific Patterns

GLM-5 (58.1%): co-leader, strongest on tool selection and the only model besides Claude to achieve high scores on ranking and province aggregation. **Claude Sonnet 4** (58.1%): co-leader, strong on municipality analysis, ranking, threshold tasks, and BigEarthNet cases (83% on the 24 BigEarthNet-related cases, highest among all models). However, it is the most expensive at \$0.087/case. **DeepSeek V3.2** (52.7%): balanced across categories with particular strength in spatial reasoning, memory/multi-turn, and municipality analysis. **Qwen3-235B** (47.3%): strongest on error handling, single analysis, and habitat analysis, but weakest on tool selection—correctly reasoning about *what* to do but failing to invoke the right tool. **Gemini 2.5 Pro** (39.8%): generated plausible-sounding answers from training data *without calling any tools* in many cases—a hallucination pattern where the model bypasses the tool-calling protocol entirely. **GPT-OSS-120B** (39.8%): moderate performance across categories with no standout strengths. **Llama 4 Scout** (5.4%): only succeeded on a handful of cases (municipality and memory), failing all other categories; consistent with Vertex MaaS infrastructure instability for Meta models.

4.5 Category-Level Analysis

Figure 8 shows accuracy by category across the six completed models. Three key findings emerge: (1) **comparison is universally unsolved** (0% across all models); (2) **er-**




Figure 7: Universally hard categories—average accuracy across all completed models.




Figure 8: Accuracy by category for six completed models. Comparison achieves 0% across all models; error handling is near-solved.

ror handling is near-solved (17–100%, with Qwen3 achieving 100%); (3) **language support is consistently strong** for the top six models on Galician and Basque inputs. The 21-case tool selection category—the largest—shows high variance across models, revealing that correct tool identification is model-dependent rather than category-dependent.

5 Discussion

5.1 Comparison with Prior Benchmarks

GeoAgentBench’s 58.1% best-model accuracy is notably lower than GeoBenchX and GeoJSON Agents (85–97%). Table 7 summarises the key differences.

Four factors explain the accuracy gap: (1) real API interaction vs local file operations; (2) environmental domain complexity (geographic knowledge, data limitations, display-only layers); (3) multi-turn conversation history; (4) strict binary evaluation. Our results align with harder benchmarks: MapEval (<67%) [Dihan et al. \[2025\]](#), SpatiaLab (55%) [Wasi et al. \[2026\]](#), UnivEARTH (33%) [Kao et al. \[2025\]](#).

The GeoBenchX comparison is particularly instructive

Table 7: Comparison of structured tool-calling geospatial benchmarks.

	GeoBenchX	GeoAgentBench
Tasks	202	93
Categories	5	18
Tools	24 primitives	12 domain ops
Tool abstraction	Low-level	High-level
Data source	Local files	Cloud API (COG + JSON)
Domain	General GIS	Environmental
Countries	—	Spain + Portugal
Indicators	—	3
Multi-turn	No	Yes
Best accuracy	85–90%	58.1%
Evaluation	LLM-as-judge	Metric-based

because both benchmarks share the structured tool-calling paradigm. The 27–39 percentage-point accuracy drop from general GIS tasks to environmental analysis suggests that general-purpose geospatial benchmarks substantially overestimate agent capability for domain-specific workflows. Environmental tasks require geographic knowledge (e.g., which provinces have erosion data), awareness of data limitations (display-only layers that cannot be analysed), multi-step reasoning across indicators, and temporal awareness—challenges absent from generic filter-merge-visualise pipelines. This motivates domain-specific benchmarks as a necessary complement to general-purpose evaluations.

5.2 Structured Tool Calling vs Code Generation

Our 58.1% on real API tool calls compares favourably with UnivEARTH’s 33% on code generation [Kao et al. \[2025\]](#). While not directly comparable, the structured approach constrains the output space, eliminating the 58% code execution failure rate. This provides further evidence that production environmental systems should prefer structured APIs over free-form code generation.

5.3 Cross-Dataset Evaluation: BigEarthNet V2

GeoAgentBench is designed to be domain- and geography-agnostic. To demonstrate extensibility beyond Spanish CO2 and erosion indicators, we incorporate BigEarthNet V2 [Clasen et al. \[2024\]](#) as a third data domain covering a different country. BigEarthNet V2 provides 549k labeled Sentinel-2 patches with CLC2018 19-class labels; we aggregate the 75k+ Portuguese patches across 9 mainland districts into a 7-class land cover distribution (Coniferous Forest, Broadleaf Forest, Shrubland, Grassland, Sparse Vegetation, Water, Urban) queryable via the same tool used for Spanish indicators.

Table 8: BigEarthNet V2 benchmark cases by task pattern. All cases use Portuguese districts with pre-computed ground truth.

Task Pattern	Cases
Single-district land cover breakdown	3
Cross-indicator (PT land cover + ES)	3
Multi-turn recall from session history	2
Chart generation	2
Legend / list layers discovery	2
Deep-dive multi-indicator profile	4
Two-district comparison	4

This integration tests whether agents generalise from Spanish environmental data to a different country and data domain. The 20 BigEarthNet benchmark cases (Table 8) span seven task patterns: single-district breakdown, cross-indicator synthesis (Portuguese land cover with Spanish CO₂/erosion), multi-turn recall, chart generation, legend and layer discovery, temporal context, and two-district comparison. Unlike the CO₂ and erosion indicators where ground truth depends on live raster geometry, BigEarthNet cases have fixed numeric ground truth derived from pre-computed district-level statistics, enabling deterministic scoring.

While Table 8 lists the 20 purpose-built BigEarthNet cases, Portuguese land cover data also appears in 4 additional cases across other categories (tool selection, interpretation, error recovery), yielding 24 BigEarthNet-related cases in total. Across all 24, Claude Sonnet 4 achieves the highest accuracy (83.3%), followed by GLM-5 (58.3%), and DeepSeek V3.2, Qwen3-235B, and Gemini 2.5 Pro (all 54.2%), with GPT-OSS-120B at 33.3%. Claude’s dominance on the BigEarthNet cases—despite sharing overall first place with GLM-5 at 58.1%—suggests that its native API integration and stronger instruction-following enable more reliable handling of the Portuguese district lookup and JSON-backed indicator workflow. The three models tied at 54.2% demonstrate consistent cross-country generalisation, while the gap to GPT-OSS-120B (33.3%) highlights model-specific weaknesses in multi-step tool orchestration.

The inclusion of BigEarthNet V2 as a queryable layer rather than merely a validation dataset transforms the benchmark’s geographic scope from Spain-only to cross-country (Spain + Portugal), and its indicator count from two to three. This validates the claim that the framework generalises to new data sources: adding a new indicator required only a JSON statistics file, a prompt update, and new benchmark cases—no changes to the scoring engine or evaluation protocol.

5.4 Limitations

Geographic scope: benchmark tasks target Spanish provinces/municipalities and Portuguese districts, covering the Iberian Peninsula but not yet extending to other re-

gions. **Indicators:** three analyzable layers (CO₂ suitability, gully erosion, BigEarthNet V2 land cover), though 18 task categories provide broad behavioural coverage. **Infrastructure availability:** Llama 4 Maverick returned persistent HTTP 500 errors from Vertex AI across all 93 cases, yielding 0% accuracy that reflects infrastructure failure rather than model capability. **Single-run:** temperature 1.0 introduces variance; no multi-run confidence intervals. **No prompt ablation:** all models use v3 zero-shot. **Evaluation strictness:** binary pass/fail may understate capability (Section 3.3 describes the partial-credit metrics that address this).

6 Conclusion

We introduce GeoAgentBench, a benchmark for evaluating AI agents on environmental geospatial analysis through structured tool calling against a production API. The benchmark spans three indicators across Spain and Portugal, demonstrating domain and geographic extensibility. Evaluating eight LLMs on 93 tasks across 18 categories with twelve tools:

1. **Environmental geospatial tool orchestration remains open.** The best models (GLM-5 and Claude Sonnet 4) achieve 58.1%—well below the 85–97% on general GIS benchmarks, with only three models exceeding 50%.
2. **Cost-accuracy trade-offs are dramatic.** DeepSeek V3.2 achieves 52.7% at \$0.008/case; Claude Sonnet 4 costs 11× more for the highest accuracy (58.1%).
3. **Systematic failures persist.** No model passes comparison tasks; deep-dive profiling averages 30% across models.
4. **Real API interaction is more discriminative** than synthetic benchmarks, with a 27–32 percentage-point gap versus general GIS evaluations.

Future work: expanding to additional countries and indicators (NDVI, precipitation, biodiversity indices); prompt ablation (zero-shot, CoT, few-shot); multi-run evaluation with confidence intervals; multi-agent architecture comparison following the Planner-Worker paradigm of GeoJSON Agents Luo et al. [2025].

GeoAgentBench, evaluation code, and the full MLOps pipeline are publicly available.²

Acknowledgments

The authors thank Darwins Geospatial S.L. for providing the production API infrastructure used in this evaluation.

²

References

- T. Akinboyewa et al. GIS Copilot: Towards an autonomous GIS agent for spatial analysis. *arXiv preprint*, 2024.
- C.M. Arnold et al. MapQA: A map-question-answering benchmark for visual language model reasoning. *Submitted to ICLR 2026*, 2025. OpenReview:dOISCbmkmg.
- Kai Norman Clasen, Leonard Hackel, Tom Burgert, Gencer Sumbul, Begüm Demir, and Volker Markl. reBEN: Refined BigEarthNet dataset for remote sensing image analysis. *arXiv preprint*, 2024.
- M.L. Dihan et al. MapEval: A map-based evaluation of geo-spatial reasoning in foundation models. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2025.
- Google Research. Google Earth AI: Unlocking geospatial insights with foundation models and cross-modal reasoning. Google Research Blog, 2025.
- Y. Huang et al. EnviroExam: Benchmarking environmental science knowledge of large language models. *arXiv preprint*, 2024.
- C.H. Kao et al. Towards LLM agents for earth observation. *arXiv preprint*, 2025.
- Kingdom of Spain. Real decreto 214/2025 on the regulation of the carbon footprint register. Boletín Oficial del Estado, April 2025.
- V. Krechetova et al. GeoBenchX: Benchmarking LLMs in agent solving multistep geospatial tasks. *arXiv preprint*, 2025.
- A. Lacoste et al. GEO-Bench: Toward foundation models for earth monitoring. *arXiv preprint*, 2023.
- C. Lee et al. Multi-agent geospatial copilots for remote sensing workflows. *arXiv preprint*, 2025.
- Z. Li and H. Ning. Autonomous GIS: the next-generation AI-powered GIS. *arXiv preprint*, 2023.
- Q. Luo et al. GeoJSON Agents: A multi-agent LLM architecture for geospatial analysis—function calling vs code generation. *arXiv preprint*, 2025.
- V. Marsocci et al. PANGAEA: A global and inclusive benchmark for geospatial foundation models. *arXiv preprint*, 2024.
- Y.-E. Pan et al. LLM-enhanced air quality monitoring interface via model context protocol. *arXiv preprint*, 2025.
- Planet Labs and Anthropic. Planet and anthropic partner to use claude’s advanced AI capabilities to turn geospatial satellite imagery into actionable insights. BusinessWire, March 2025.
- A. Shabbir et al. ThinkGeo: Evaluating tool-augmented agents for remote sensing tasks. *arXiv preprint*, 2025.
- N. Simumba et al. GEO-Bench-2: From performance to capability, rethinking evaluation in geospatial AI. *arXiv preprint*, 2025.
- D. Stamoulis et al. Geo-OLM: Enabling sustainable earth observation studies with cost-efficient open language models & state-driven workflows. *arXiv preprint*, 2025.
- D. Szwarcman et al. Prithvi-EO-2.0: A versatile multi-temporal foundation model for earth observation applications. *arXiv preprint*, 2024.
- T.H. Truong et al. GPSBench: Do large language models understand GPS coordinates? *arXiv preprint*, 2026.
- A.T. Wasi et al. SpatiaLab: Can vision-language models perform spatial reasoning in the wild? *arXiv preprint*, 2026.
- S. Yao et al. ReAct: Synergizing reasoning and acting in language models. *arXiv preprint*, 2023.
- S. Zhang et al. GeoGPT: Understanding and processing geospatial tasks through an autonomous GPT. *arXiv preprint*, 2023.
- Y. Zhang et al. Geospatial large language model trained with a simulated environment for generating tool-use chains autonomously. *International Journal of Applied Earth Observation and Geoinformation*, 2024. doi: 10.1016/j.jag.2024.104295.
- Y. Zhang et al. Fine-tuning large language models for interdisciplinary environmental challenges. *Environmental Science and Ecotechnology*, 2025. doi: 10.1016/j.ese.2025.100608.